

```
!pip install ucimlrepo

Collecting ucimlrepo
  Downloading ucimlrepo-0.0.6-py3-none-any.whl (8.0 kB)
Installing collected packages: ucimlrepo
Successfully installed ucimlrepo-0.0.6

from ucimlrepo import fetch_ucirepo

# fetch dataset
census_income = fetch_ucirepo(id=20)

# data (as pandas dataframes)
X = census_income.data.features
y = census_income.data.targets

# metadata
print(census_income.metadata)

# variable information
print(census_income.variables)

{'uci_id': 20, 'name': 'Census Income', 'repository_url': 'https://archive.ics.uci.edu/dataset/20/census+income', 'data_url': 'https://archive.ics.uci.edu/dataset/20/census+income'}
name      role      type      demographic \
0         age  Feature    Integer          Age
1    workclass  Feature  Categorical      Income
2      fnlwgt  Feature    Integer          None
3    education  Feature  Categorical  Education Level
4  education-num  Feature    Integer  Education Level
5  marital-status  Feature  Categorical      Other
6    occupation  Feature  Categorical      Other
7  relationship  Feature  Categorical      Other
8         race  Feature  Categorical      Race
9         sex  Feature    Binary      Sex
10  capital-gain  Feature    Integer      None
11  capital-loss  Feature    Integer      None
12  hours-per-week  Feature    Integer      None
13 native-country  Feature  Categorical      Other
14      income  Target      Binary      Income

description units missing_values
0              N/A  None          no
1  Private, Self-emp-not-inc, Self-emp-inc, Feder...  None          yes
2              None  None          no
3  Bachelors, Some-college, 11th, HS-grad, Prof-...  None          no
4              None  None          no
5  Married-civ-spouse, Divorced, Never-married, S...  None          no
6  Tech-support, Craft-repair, Other-service, Sal...  None          yes
7  Wife, Own-child, Husband, Not-in-family, Other...  None          no
8  White, Asian-Pac-Islander, Amer-Indian-Eskimo,...  None          no
9              Female, Male.  None          no
10             None  None          no
11             None  None          no
12             None  None          no
13  United-States, Cambodia, England, Puerto-Rico,...  None          yes
14             >50K, <=50K.  None          no
```

```
import pandas as pd
import numpy as np
```

X

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husbanc
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husbanc
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife
...

y

	income
0	<=50K
1	<=50K
2	<=50K
3	<=50K
4	<=50K
...	...
48837	<=50K.
48838	<=50K.
48839	<=50K.
48840	<=50K.
48841	>50K.

48842 rows × 1 columns

```
merged_census = pd.concat([X, y], axis = 1)
merged_census
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husbanc
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husbanc
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife
...

```
# checks for duplicates
merged_census[merged_census.duplicated()].shape[0]

29

# removes duplicate rows
merged_census.drop_duplicates(inplace=True)

merged_census.info()
# it can be noticed that workclass, occupation, and native-country has missing data

<class 'pandas.core.frame.DataFrame'>
Index: 48813 entries, 0 to 48841
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0    age                    48813 non-null  int64
1    workclass              47850 non-null  object
2    fnlwgt                 48813 non-null  int64
3    education              48813 non-null  object
4    education-num          48813 non-null  int64
5    marital-status         48813 non-null  object
6    occupation             47847 non-null  object
7    relationship           48813 non-null  object
8    race                   48813 non-null  object
9    sex                    48813 non-null  object
10   capital-gain           48813 non-null  int64
11   capital-loss           48813 non-null  int64
12   hours-per-week         48813 non-null  int64
13   native-country         48539 non-null  object
14   income                 48813 non-null  object
dtypes: int64(6), object(9)
memory usage: 6.0+ MB

# working with workclass first
merged_census.workclass.unique()
# it can be noticed that there are 2 missing values (? and NaN)

array(['State-gov', 'Self-emp-not-inc', 'Private', 'Federal-gov',
       'Local-gov', '?', 'Self-emp-inc', 'Without-pay', 'Never-worked',
       nan], dtype=object)

# checking the frequency of the unique values
merged_census.workclass.value_counts()
# as Private is the most frequent, it would be more appropriate
# to change the missing values to Private in order for there to be
# less significant skewing of data

workclass
Private      33879
```

```
Self-emp-not-inc    3861
Local-gov          3136
State-gov          1981
?                  1836
Self-emp-inc       1694
Federal-gov        1432
Without-pay        21
Never-worked       10
Name: count, dtype: int64
```

```
merged_census.workclass.fillna('Private', inplace=True)
merged_census.workclass.replace('?', 'Private', inplace=True)
merged_census
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husbanc
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husbanc
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife
...

```
merged_census.workclass.unique()

array(['State-gov', 'Self-emp-not-inc', 'Private', 'Federal-gov',
      'Local-gov', 'Self-emp-inc', 'Without-pay', 'Never-worked'],
      dtype=object)

# working with occupation
merged_census.occupation.unique()

array(['Adm-clerical', 'Exec-managerial', 'Handlers-cleaners',
      'Prof-specialty', 'Other-service', 'Sales', 'Craft-repair',
      'Transport-moving', 'Farming-fishing', 'Machine-op-inspct',
      'Tech-support', '?', 'Protective-serv', 'Armed-Forces',
      'Priv-house-serv', nan], dtype=object)

merged_census.occupation.value_counts()
# as the data is more evenly distributed, I decided to use forward fill
```

```
occupation
Prof-specialty    6167
Craft-repair      6107
Exec-managerial   6084
Adm-clerical      5608
Sales             5504
Other-service     4919
Machine-op-inspct 3019
Transport-moving  2355
Handlers-cleaners 2071
?                1843
Farming-fishing   1487
Tech-support      1445
Protective-serv   983
Priv-house-serv   240
Armed-Forces      15
Name: count, dtype: int64
```

```
merged_census.occupation.fillna(method='ffill', inplace=True)
```

```
merged_census.occupation.replace('?', method='ffill', inplace=True)
merged_census
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husbanc
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husbanc
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife
...

```
# working with native-country
merged_census['native-country'].unique()

array(['United-States', 'Cuba', 'Jamaica', 'India', '?', 'Mexico',
      'South', 'Puerto-Rico', 'Honduras', 'England', 'Canada', 'Germany',
      'Iran', 'Philippines', 'Italy', 'Poland', 'Columbia', 'Cambodia',
      'Thailand', 'Ecuador', 'Laos', 'Taiwan', 'Haiti', 'Portugal',
      'Dominican-Republic', 'El-Salvador', 'France', 'Guatemala',
      'China', 'Japan', 'Yugoslavia', 'Peru',
      'Outlying-US(Guam-USVI-etc)', 'Scotland', 'Trinidad&Tobago',
      'Greece', 'Nicaragua', 'Vietnam', 'Hong', 'Ireland', 'Hungary',
      'Holand-Netherlands', nan], dtype=object)
```

```
merged_census['native-country'].value_counts()
# as United-States is the most frequent,
# it is more appropriate to change the missing values
# to the United-States
```

native-country	
United-States	43810
Mexico	947
?	582
Philippines	295
Germany	206
Puerto-Rico	184
Canada	182
El-Salvador	155
India	151
Cuba	138
England	127
China	122
South	115
Jamaica	106
Italy	105
Dominican-Republic	103
Japan	92
Poland	87
Guatemala	86
Vietnam	86
Columbia	85
Haiti	75
Portugal	67
Taiwan	65
Iran	59
Greece	49
Nicaragua	49
Peru	46
Ecuador	45
France	38
Ireland	37
Hong	30
Thailand	30
Cambodia	28
Trinidad&Tobago	27
Laos	23
Yugoslavia	23
Outlying-US(Guam-USVI-etc)	23
Scotland	21
Honduras	20
Hungary	19
Holand-Netherlands	1
Name: count, dtype: int64	

```
merged_census['native-country'].fillna('United-States', inplace=True)
merged_census['native-country'].replace('?', 'United-States', inplace=True)
merged_census
```

[illegible]

```
# it can be noticed that in income, the values are in the wrong format
merged_census.income.unique()
```

```
array(['<=50K', '>50K', '<=50K.', '>50K.'], dtype=object)
```

```
# replacing the values to be more uniform
merged_census.income.replace('<=50K.', '<=50K', inplace=True)
merged_census.income.replace('>50K.', '>50K', inplace=True)
```

```
merged_census.income.unique()

array(['<=50K', '>50K'], dtype=object)
```

```
merged_census.info()
# all missing data are handled
```

```

class 'pandas.core.frame.DataFrame'>
Index: 48813 entries, 0 to 48841
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   age                   48813 non-null  int64  
1   workclass              48813 non-null  object  
2   fnlwgt                48813 non-null  int64  
3   education              48813 non-null  object  
4   education-num         48813 non-null  int64  
5   marital-status        48813 non-null  object  
6   occupation            48813 non-null  object  
7   relationship          48813 non-null  object  
8   race                  48813 non-null  object  
9   sex                   48813 non-null  object  
10  capital-gain          48813 non-null  int64  
11  capital-loss          48813 non-null  int64  
12  hours-per-week        48813 non-null  int64  
13  native-country        48813 non-null  object  
14  income                48813 non-null  object  
dtypes: int64(6), object(9)
memory usage: 6.0+ MB

```

```
# as for now, I don't have any use for fnlwgt, capital-gain, and capital-loss,
# I decided to drop them
merged_census.drop(columns=['fnlwgt', 'capital-gain', 'capital-loss'], inplace=True)
merged_census
```

[illegible]

```
# converting the categorical data into numerical data
# as there is already education-num, which represents the education numerically
# we can already drop the education column
# but I want to have a dictionary of the corresponding values of the education

education_map = dict(zip(merged_census.education, merged_census['education-num']))

merged_census.drop(columns=['education'], inplace=True)
merged_census
```

	age	workclass	education-num	marital-status	occupation	relationship	race	sex
0	39	State-gov	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	9	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	53	Private	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	28	Private	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female
...

```
education_map

{'Bachelors': 13,
 'HS-grad': 9,
 '11th': 7,
 'Masters': 14,
 '9th': 5,
 'Some-college': 10,
 'Assoc-acdm': 12,
 'Assoc-voc': 11,
 '7th-8th': 4,
 'Doctorate': 16,
 'Prof-school': 15,
 '5th-6th': 3,
 '10th': 6,
 '1st-4th': 2,
 'Preschool': 1,
 '12th': 8}

# getting the unique values
columns = [
    'workclass', 'marital-status', 'occupation', 'relationship', 'race', 'sex', 'native-country', 'income'
] # columns to get the unique values
unique_values = []

# gets the unique values of a column and appends it to the unique_values list
for column in columns:
    unique_values.append(merged_census[column].unique().tolist())
unique_values

[['State-gov',
 'Self-emp-not-inc',
 'Private',
 'Federal-gov',
 'Local-gov',
 'Self-emp-inc',
 'Without-pay',
 'Never-worked'],
 ['Never-married',
 'Married-civ-spouse',
 'Divorced',
 'Married-spouse-absent',
 'Separated',
 'Married-AF-spouse',
 'Widowed'],
 ['Adm-clerical',
 'Exec-managerial',
 'Handlers-cleaners',
 'Prof-specialty',
 'Other-service',
 'Sales',
 'Craft-repair',
 'Transport-moving',
 'Farming-fishing',
 'Machine-op-inspct',
 'Tech-support',
 'Protective-serv',
 'Armed-Forces',
 'Priv-house-serv'],
 ['Not-in-family',
 'Husband',
 'Wife',
 'Own-child',
 'Unmarried',
 'Other-relative'],
 ['White', 'Black', 'Asian-Pac-Islander', 'Amer-Indian-Eskimo', 'Other'],
 ['Male', 'Female'],
 ['United-States',
 'Cuba',
 'Jamaica',
 'India',
```

```

    'Mexico',
    'South',
    'Puerto-Rico',
    'Honduras',
    'England',
    'Canada',
    'Germany',
    'Iran',
    'Philippines',
    'Italy',
    'Poland',
    'Columbia',
    'Cambodia',
    'Thailand',
    'Ecuador',
    'Laos',
    'Taiwan'.

# creates the dictionaries
result_dicts = [] # stores the results here

for data in unique_values:
    keys = [i for i in data]
    values = [i for i in range(1, len(data)+1)]
    result_dicts.append({keys[i] : values[i] for i in range(len(values))})
result_dicts

[{'State-gov': 1,
'Self-emp-not-inc': 2,
'Private': 3,
'Federal-gov': 4,
'Local-gov': 5,
'Self-emp-inc': 6,
'Without-pay': 7,
'Never-worked': 8},
{'Never-married': 1,
'Married-civ-spouse': 2,
'Divorced': 3,
'Married-spouse-absent': 4,
'Separated': 5,
'Married-AF-spouse': 6,
'Widowed': 7},
{'Adm-clerical': 1,
'Exec-managerial': 2,
'Handlers-cleaners': 3,
'Prof-specialty': 4,
'Other-service': 5,
'Sales': 6,
'Craft-repair': 7,
'Transport-moving': 8,
'Farming-fishing': 9,
'Machine-op-inspct': 10,
'Tech-support': 11,
'Protective-serv': 12,
'Armed-Forces': 13,
'Priv-house-serv': 14},
{'Not-in-family': 1,
'Husband': 2,
'Wife': 3,
'Own-child': 4,
'Unmarried': 5,
'Other-relative': 6},
{'White': 1,
'Black': 2,
'Asian-Pac-Islander': 3,
'Amer-Indian-Eskimo': 4,
'Other': 5},
{'Male': 1, 'Female': 2},
{'United-States': 1,
'Cuba': 2,
'Jamaica': 3,
'India': 4,
'Mexico': 5,
'South': 6,
'Puerto-Rico': 7,
'Honduras': 8,
'England': 9,
'Canada': 10,
'Germany': 11,
'Iran': 12,
'Philippines': 13,
'Italy': 14,
'Poland': 15,
'Columbia': 16,
'Cambodia': 17,

# maps the categorical data to their numerical counterparts
for column in range(len(columns)):
    merged_census.replace(result_dicts[column], inplace=True)

merged_census
```

	age	workclass	education-num	marital-status	occupation	relationship	race	sex	hours-per-week
0	39	1	13	1	1	1	1	1	40
1	50	2	13	2	2	2	1	1	13
2	38	3	9	3	3	1	1	1	40
3	53	3	7	2	3	2	2	1	40
4	28	3	13	2	4	3	2	2	40
...
48837	39	3	13	3	4	1	1	2	36
48838	64	3	9	7	4	6	2	1	40
48839	38	3	13	2	4	2	1	1	50
48840	44	3	13	3	1	4	3	1	40
48841	35	6	13	2	2	2	1	1	60

```
merged_census.describe()
```

	age	workclass	education-num	marital-status	occupation	relationship
count	48813.000000	48813.000000	48813.000000	48813.000000	48813.000000	48813.000000
mean	38.647348	3.104419	10.078688	2.084322	5.276484	2.539037
std	13.709005	0.917171	2.570257	1.257648	3.044477	1.440102
min	17.000000	1.000000	1.000000	1.000000	1.000000	1.000000
25%	28.000000	3.000000	9.000000	1.000000	2.000000	1.000000
50%	37.000000	3.000000	10.000000	2.000000	5.000000	2.000000
75%	48.000000	3.000000	12.000000	2.000000	7.000000	4.000000

```
# comparing easy to distinguish categories (sex and income)
male_census = merged_census.query('sex == 1')
female_census = merged_census.query('sex == 2')
less_than_census = merged_census.query('income == 1')
more_than_census = merged_census.query('income == 2')
```

```
male_census.mean()

age          39.497594
workclass    3.110570
education-num 10.095492
marital-status 1.928320
occupation   5.687383
relationship  2.262389
race         1.191107
sex          1.000000
hours-per-week 42.419264
native-country 2.118262
income       1.303883
dtype: float64
```

```
female_census.mean()

age          36.932827
workclass    3.092016
education-num 10.044803
marital-status 2.398900
occupation   4.447905
relationship  3.096898
race         1.279137
sex          2.000000
hours-per-week 36.403720
native-country 2.141824
income       1.109319
dtype: float64
```

```
# on average, men are reported to be older than women
# men have also received slightly higher education than women
# men also work longer per week than women
# and men earn more than women
# women on the other hand are in their later stage of relationships compared to men
```

```
less_than_census.mean()

age          36.875916
workclass    3.053787
education-num 9.598901
marital-status 2.082660
occupation   5.412734
relationship  2.677063
race         1.242512
sex          1.388198
hours-per-week 38.842599
native-country 2.191419
income       1.000000
dtype: float64
```



```
more_than_census.mean()


age          44.275909
workclass    3.265297
education-num 11.603166
marital-status 2.089602
occupation   4.843560
relationship  2.100471
race         1.149679
sex          1.151391
hours-per-week 45.453145
native-country 1.918442
income       2.000000
dtype: float64

# on average, people who earn more are older
# they also have higher educational attainment
# they also work longer hours
```

Plotting

```
%matplotlib inline
import matplotlib.pyplot as plt

import seaborn as sns
plt.figure(figsize=(10,10))
sns.heatmap(
    merged_census.sort_index().corr(),
    annot=True, center=0, square=True
)
# it can be noticed that there are notable correlations
# with marital status and age, education and income,
# race and native country, and hours per week and income
```



```
# visualizing the relationship between age and marital status


# creates subplots
fig, axes = plt.subplots(figsize=(8, 8))
# gets the different marital status
statuses = merged_census['marital-status'].unique()

# creates a histogram which per marital statis
for status in statuses:
    sns.kdeplot(merged_census[merged_census['marital-status'] == status]['age'],
        ax=axes, label=status)

plt.suptitle('KDE of Age by Marital Status')
plt.xlabel('Age')
plt.legend(unique_values[1], title='Marital Status')


plt.savefig('age_marital-status.png')
```

KDE of Age by Marital Status



```
# as both age and hours-per-week are the only datapoints left that are non-categorical,
# I decided to graph them
sns.jointplot(
    x='age',
    y='hours-per-week',
    kind='hex', # hex plot
    data=merged_census
)
# it can then be noticed that there is a lot of people working 40 hours per week

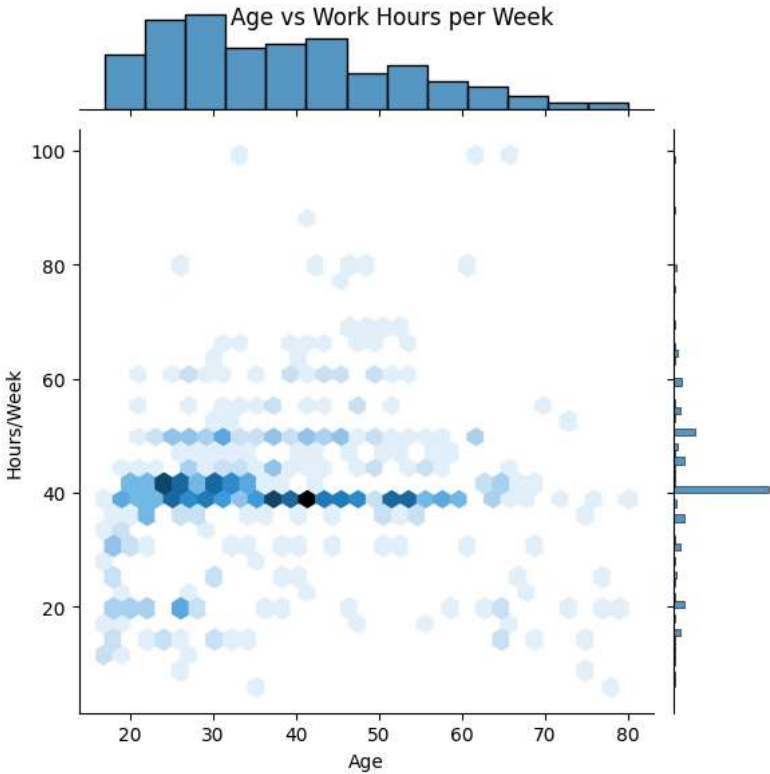
<seaborn.axisgrid.JointGrid at 0x7e3425d5eef0>
```



```
# however, the data looks too small,
# therefore I decided to get a sample of 500
census_sample = merged_census.sample(n=500, random_state=0)
census_sample
```

	age	workclass	education-num	marital-status	occupation	relationship	race	sex	hours-per-week
2293	52	3	10	2	7	2	1	1	40
7348	25	1	10	2	5	3	1	2	15
40927	28	3	9	2	7	2	1	1	40
22311	19	3	9	1	5	4	2	2	40
44795	29	3	4	2	7	2	1	1	40
...
24803	18	3	9	1	1	4	1	2	20
13739	21	3	10	1	6	4	2	2	20
43071	32	2	9	2	7	2	1	1	32
28013	19	3	10	1	2	4	1	2	40
1433	18	3	10	1	6	4	1	2	15

```
sns.jointplot(
    x='age',
    y='hours-per-week',
    kind='hex', # hex plot
    data=census_sample,
)
plt.xlabel('Age')
plt.ylabel('Hours/Week')
plt.suptitle('Age vs Work Hours per Week')
# it can be seen that most of the data is at around
# age 20-50 and are working 20-60 hours per week
plt.savefig('age_hours-week.png')
```



```
# looking at categorical data
# as sex and income are both binary, I decided to use them for plotting
a = male_census.agg({
    'age' : 'mean',
    'education-num' : 'mean',
    'hours-per-week' : 'mean',
    'income' : 'mean',
    'marital-status' : 'mean'
})
a

age          39.497594
education-num 10.095492
hours-per-week 42.419264
income        1.303883
marital-status 1.928320
dtype: float64

b = female_census.agg({
    'age' : 'mean',
    'education-num' : 'mean',
    'hours-per-week' : 'mean',
    'income' : 'mean',
    'marital-status' : 'mean'
})
b

age          36.932827
education-num 10.044803
hours-per-week 36.403720
income        1.109319
```

```
marital-status      2.398900
dtype: float64

by_sex_census = pd.concat([a, b], axis = 1)
by_sex_census
```

	0	1
age	39.497594	36.932827
education-num	10.095492	10.044803
hours-per-week	42.419264	36.403720
income	1.303883	1.109319
marital-status	1.928320	2.398900

```
by_sex_census.rename(columns={0 : 'male', 1 : 'female'}, inplace=True)
by_sex_census.index=['Age', 'Education', 'Hours/Week', 'Income', 'Marital Status']
by_sex_census
```

	male	female
Age	39.497594	36.932827
Education	10.095492	10.044803
Hours/Week	42.419264	36.403720
Income	1.303883	1.109319
Marital Status	1.928320	2.398900


```
# creating the bar plot
low_values = by_sex_census.iloc[[1,3,4]]
high_values = by_sex_census.iloc[[0,2]]

fig, (ax_low, ax_high) = plt.subplots(1, 2, figsize=(12, 4))

low_values.plot(
    kind='bar',
    cmap='tab20b',
    ax=ax_low
)

high_values.plot(
    kind='bar',
    cmap='tab20b',
    ax=ax_high
)
fig.suptitle('Sex Averages')

plt.savefig('sex_average.png')
```



```
c = less_than_census.agg({
    'age' : 'mean',
    'education-num' : 'mean',
    'hours-per-week' : 'mean',
})
c

age      36.875916
education-num  9.598901
hours-per-week 38.842599
dtype: float64
```

```
d = more_than_census.agg({
    'age' : 'mean',
    'education-num' : 'mean',
    'hours-per-week' : 'mean',
})
d
```

```
age          44.275909
education-num 11.603166
hours-per-week 45.453145
dtype: float64
```

```
by_income_census = pd.concat([c, d], axis = 1)
by_income_census
```

	0	1
age	36.875916	44.275909
education-num	9.598901	11.603166
hours-per-week	38.842599	45.453145

```
by_income_census.rename(columns={0 : '<=50K', 1 : '>50K'}, inplace=True)
by_income_census.index = ['Age', 'Education', 'Hours/Week']
by_income_census
```

	<=50K	>50K
Age	36.875916	44.275909
Education	9.598901	11.603166
Hours/Week	38.842599	45.453145

```
# creating the bar plot
by_income_census.plot(
    kind='bar',
    cmap = 'tab20c',
    title = 'Income Averages'
)
plt.savefig('income_average.png')
```

